Noa Shapira

Haifa, Israel · noa.shapira@example.com · Open to remote and hybrid

Summary

Machine learning engineer with five years of experience, focused on taking models past the notebook and into systems people actually depend on. Half my work is modelling and half is the evaluation, monitoring and data plumbing that decides whether the modelling was worth anything. Recently spent a lot of time on LLM applications and their evaluation.

Experience

Senior ML Engineer — Refua Health AI, Haifa · Apr 2023 – present

- Own the clinical note summarization product end to end: retrieval, prompting, evaluation harness and the human review queue used by twelve partner clinics.
- Built an LLM-as-judge evaluation suite over 1,200 clinician-labelled examples; it caught a regression that would otherwise have shipped a 9-point drop in factuality.
- Reduced inference cost per note by 62% through prompt caching, aggressive retrieval pruning, and moving the easy cases to a smaller model.
- Work directly with clinicians on failure review — the single highest-value hour in my week and the source of most real model improvements.

Data Scientist — Migdal Analytics, Tel Aviv · Feb 2021 – Mar 2023

- Built churn and lifetime-value models for subscription clients in PyTorch and scikit-learn; the churn model drove a retention campaign worth about $1.4M annually.
- Replaced a fragile collection of notebooks with an Airflow and dbt pipeline, which made the weekly client reports reproducible for the first time.
- Ran the A/B testing framework and, more often than anyone liked, explained why a result was not significant.

Skills

Languages: Python, SQL, some Go
ML: PyTorch, scikit-learn, Hugging Face Transformers, XGBoost
LLM work: RAG, evaluation harnesses, prompt engineering, vector search (pgvector, Pinecone)
Data and platform: Airflow, dbt, Spark, Snowflake, Docker, AWS SageMaker, Weights & Biases

Education

M.Sc. Computer Science (machine learning), Technion, 2021. Thesis on uncertainty estimation for medical image classifiers. B.Sc. Statistics, Hebrew University, 2019.

Publications

Second author on a workshop paper about calibration under distribution shift (NeurIPS workshop, 2022). Write an occasional newsletter on LLM evaluation with ~2,000 readers.

Languages: Hebrew (native), English (fluent)